

ANSHU AGARWAL

DATA ENGINEER

+91 90249 45394
anshuagarwal979@gmail.com
Pune, Maharashtra
linkedin.com/in/anshuagarwal14
github.com/anshu1425

PROFESSIONAL SUMMARY

GCP Data Engineer with **3+ years** turning fragile, manual data operations into resilient, scalable data platforms and automated pipelines for global enterprises like **General Mills and Bayer**. Delivered multi-terabyte cloud migrations, improved **ETL performance by 25%**, and strengthened data reliability through observability and data quality frameworks. Skilled in **BigQuery, Apache Airflow, dbt, Python, SQL, and Spark**, while leveraging **GenAI-assisted engineering** practices to accelerate delivery and streamline engineering workflows. Recognized for end-to-end ownership, rapid ramp-up, and measurable reliability and efficiency gains.

WORK EXPERIENCE

Associate Data Engineer | COGNIZANT TECHNOLOGY SOLUTIONS

March 2022 – Present

Project: NAF Data Foundation – General Mills Inc.

(Current)

- **Data Architecture & Ingestion:** Built production-grade ETL pipelines using **Apache Airflow, dbt, and BigQuery** to establish a unified canonical data layer powering **predictive analytics** for enterprise sales and supply-chain operations.
- **Data Pipeline Observability:** Engineered an automated monitoring system using secure notification protocols for critical production DAGs, reducing alert response times from **~45 minutes to under 5 minutes** and improving **MTTR by 30-40%**.
- **Data Quality & Contracts:** Deployed comprehensive **unit and generic dbt test suites** to validate business schemas, data types, and integrity constraints; operationalized **8 fail-fast dbt Cloud jobs** inside production DAGs to block **~200-300 duplicate records/month** and prevent schema drift while expanding detection coverage by **80%**.
- **Data Schema & Pipeline Optimization:** Optimized high-volume supply-chain ingestion with dynamic file parsing and schema-deviation logic; reduced data processing time from over a week to under 24 hours and slashed month-end cycles by **85%**.
- **Master Data Governance:** Deployed format validation constraints across enterprise pipelines to normalize **50K-70K identity attributes**, guaranteeing reporting consistency between **CRM systems** and analytical downstream BI layers.
- **AI-Assisted Data Engineering:** Operationalized custom AI agents within **Claude** across **requirement analysis and deployment planning**; compressed new distributor onboarding cycles from **~7 weeks to 2-3 weeks** and reduced documentation tracking overheads by **90%**.

Project: PET ITQ-PLM Data Migration – General Mills Inc.

Jun 2024 – Jan 2026

- **Cloud Data Migration:** Architected scalable GCP-based ETL tracks for the zero-loss migration of **6+ SAP PLM datasets (5+ TB)** to Optiva, reducing manual processing latency by **25%** through cross-functional partnership with product owners and QA teams.
- **Dynamic Ingestion Pipelines:** Automated monthly batch processing of **200+ raw files** into **BigQuery** via parameterized **Airflow DAGs** and reusable JSON ingestion layouts.
- **dbt Modeling & Refactoring:** Modeled **59+ core dbt models** to deliver business-ready **BigQuery tables** aligned to **Optiva requirements**, increasing data-quality consistency by **15%** with automated Airflow refresh cycles and alert tracking to ensure continuous pipeline health.
- **DataOps & CI/CD Pipelines:** Integrated automated **GitHub CI/CD infrastructure workflows** for version control and rigorous code reviews, accelerating deployment execution velocity while sustaining a **99.9% production uptime record**.

Project: Bayer Cyrus CPH-CDP – Healthcare Data Platform – Bayer A.G.

Mar 2023 – May 2024

- **Cloud Infrastructure & SLAs:** Managed and optimized **AWS Glue ETL jobs** for high-volume batch data loading and transformations, partnering with QA and Engineering teams to deliver change requests (**CRs**) under strict **SLA adherence**.
- **Query Performance Optimization:** Conducted root-cause analysis on production load failures to restore platform health; tuned **Snowflake analytical queries** to maximize data retrieval efficiency and triage system incidents.

Internship: Full-Stack Data Integration Intern

Mar 2022 – Feb 2023

- Built foundational data-integration skills in SQL, Python, Unix shell scripting, and DBMS.
- Developed a hands-on project using enterprise ETL tools including Informatica, Talend, and DataStage.

TECHNICAL SKILLS

Data Warehousing: BigQuery, Databricks, MySQL, Power BI, Dimensional Modeling (Star/Snowflake), Data Lake

Data Pipelines: Apache Airflow, Cloud Composer, dbt Cloud, AWS Glue, ETL/ELT, MOVEit

Programming: Python, PySpark, SQL, Unix Shell Scripting

Cloud Platforms: GCP (BigQuery, GCS), AWS (S3, EC2, DynamoDB, CloudWatch), Apache Spark

Data Quality: dbt Tests, Data Standardization, Schema Validation, Deduplication

DevOps: GitHub, GitHub Actions, Azure DevOps, Agile/Scrum, Confluence

AI Tooling: Claude (Certified Architect), GitHub Copilot (Certified), Glean Agentic Workflows

Other Skills: Problem-Solving, Team Collaboration, Communication, Analytical Thinking

CERTIFICATIONS

- **Claude Certified Architect: Foundations** – Anthropic (2026)
- **Astronomer Certification for Apache Airflow 3 Fundamentals** – Astronomer (2025)
- **GitHub Copilot Certification** – Microsoft (2025)
- **The Complete SQL Bootcamp: Zero to Hero** – Udemy (2024)
- **Microsoft Certified: Power BI Data Analyst Associate (PL-300)** – Microsoft (2024)

RECOGNITIONS

- **Cognizant CHEERS** – "Promising New Star" Team Award (2024)

EDUCATION

Bachelor of Technology in Electrical Engineering (90.2%) – Poornima College of Engineering, Jaipur

2022